

Recommendations_with_IBM

May 25, 2020

1 Recommendations with IBM

In this notebook, you will be putting your recommendation skills to use on real data from the IBM Watson Studio platform.

You may either submit your notebook through the workspace here, or you may work from your local machine and submit through the next page. Either way assure that your code passes the project [RUBRIC](#). **Please save regularly.**

By following the table of contents, you will build out a number of different methods for making recommendations that can be used for different situations.

1.1 Table of Contents

I. Section ?? II. Section ?? III. Section ?? IV. Section ?? V. Section ?? VI. Section ??

At the end of the notebook, you will find directions for how to submit your work. Let's get started by importing the necessary libraries and reading in the data.

```
In [62]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import project_tests as t
import pickle

%matplotlib inline

df = pd.read_csv('data/user-item-interactions.csv')
df_content = pd.read_csv('data/articles_community.csv')
del df['Unnamed: 0']
del df_content['Unnamed: 0']

# Show df to get an idea of the data
df.head()
```

Out[62]:

	article_id	title \
0	1430.0	using pixiedust for fast, flexible, and easier...
1	1314.0	healthcare python streaming application demo
2	1429.0	use deep learning for image classification
3	1338.0	ml optimization using cognitive assistant
4	1276.0	deploy your python model as a restful api

```

                                email
0  ef5f11f77ba020cd36e1105a00ab868bbdbf7fe7
1  083cbdfa93c8444beaa4c5f5e0f5f9198e4f9e0b
2  b96a4f2e92d8572034b1e9b28f9ac673765cd074
3  06485706b34a5c9bf2a0ecdac41daf7e7654ceb7
4  f01220c46fc92c6e6b161b1849de11faacd7ccb2

```

```

In [63]: # Show df_content to get an idea of the data
df_content.head()

```

```

Out[63]:                                doc_body \
0  Skip navigation Sign in SearchLoading...\r\n\r...
1  No Free Hunch Navigation * kaggle.com\r\n\r\n ...
2  * Login\r\n * Sign Up\r\n\r\n * Learning Pat...
3  DATALAYER: HIGH THROUGHPUT, LOW LATENCY AT SCA...
4  Skip navigation Sign in SearchLoading...\r\n\r...

```

```

                                doc_description \
0  Detect bad readings in real time using Python ...
1  See the forest, see the trees. Here lies the c...
2  Heres this weeks news in Data Science and Bi...
3  Learn how distributed DBs solve the problem of...
4  This video demonstrates the power of IBM DataS...

```

	doc_full_name	doc_status	article_id
0	Detect Malfunctioning IoT Sensors with Streami...	Live	0
1	Communicating data science: A guide to present...	Live	1
2	This Week in Data Science (April 18, 2017)	Live	2
3	DataLayer Conference: Boost the performance of...	Live	3
4	Analyze NY Restaurant data using Spark in DSX	Live	4

1.1.1 Part I: Exploratory Data Analysis

Use the dictionary and cells below to provide some insight into the descriptive statistics of the data.

1. What is the distribution of how many articles a user interacts with in the dataset? Provide a visual and descriptive statistics to assist with giving a look at the number of times each user interacts with an article.

```


In [64]: interactions_count = df.groupby('email')['article_id'].count().values
plt.hist(interactions_count, bins=20)
plt.title('Number of user articles interactions')
plt.xlabel('Number of articles interactions')
plt.ylabel('Number of articles')

```

```

Out[64]: Text(0,0.5,'Number of articles')

```



```
In [65]: df.groupby('email')['article_id'].count().describe()
```

```
Out[65]: count      5148.000000
         mean         8.930847
         std        16.802267
         min         1.000000
         25%         1.000000
         50%         3.000000
         75%         9.000000
         max        364.000000
         Name: article_id, dtype: float64
```

```
In [66]: # Fill in the median and maximum number of user_article interactios below
```

```
median_val = 3 # 50% of individuals interact with ___ number of articles or fewer.
max_views_by_user = 364 # The maximum number of user-article interactions by any 1 user
```

2. Explore and remove duplicate articles from the **df_content** dataframe.

```
In [67]: # Find and explore duplicate articles
         df_content[df_content.duplicated(['article_id'])]
```

```
Out[67]: doc_body \
365 Follow Sign in / Sign up Home About Insight Da...
```

```

692 Homepage Follow Sign in / Sign up Homepage * H...
761 Homepage Follow Sign in Get started Homepage *...
970 This video shows you how to construct queries ...
971 Homepage Follow Sign in Get started * Home\r\n...

```

```

doc_description \
365 During the seven-week Insight Data Engineering...
692 One of the earliest documented catalogs was co...
761 Todays world of data science leverages data f...
970 This video shows you how to construct queries ...
971 If you are like most data scientists, you are ...

```

	doc_full_name	doc_status	article_id
365	Graph-based machine learning	Live	50
692	How smart catalogs can turn the big data flood...	Live	221
761	Using Apache Spark as a parallel processing fr...	Live	398
970	Use the Primary Index	Live	577
971	Self-service data preparation with IBM Data Re...	Live	232

```

In [68]: # example of duplicated ID
df_content[df_content.article_id == 50]

```

```

Out[68]: doc_body \
50 Follow Sign in / Sign up Home About Insight Da...
365 Follow Sign in / Sign up Home About Insight Da...

```

```

doc_description \
50 Community Detection at Scale
365 During the seven-week Insight Data Engineering...

```

	doc_full_name	doc_status	article_id
50	Graph-based machine learning	Live	50
365	Graph-based machine learning	Live	50

```

In [69]: # Remove any rows that have the same article_id - only keep the first
df_content.drop_duplicates(['article_id'], inplace=True)

```

3. Use the cells below to find:

- The number of unique articles that have an interaction with a user.
- The number of unique articles in the dataset (whether they have any interactions or not).
- The number of unique users in the dataset. (excluding null values)
- The number of user-article interactions in the dataset.

```

In [70]: df.email.nunique()

```

```

Out[70]: 5148

```

```

In [71]: len(df)

```

```

Out[71]: 45993

```

```
In [72]: unique_articles = df.article_id.nunique()# The number of unique articles that have at l
total_articles = df_content.article_id.nunique()# The number of unique articles on the
unique_users = df.email.nunique()# The number of unique users
user_article_interactions = len(df)# The number of user-article interactions
```

4. Use the cells below to find the most viewed **article_id**, as well as how often it was viewed. After talking to the company leaders, the email_mapper function was deemed a reasonable way to map users to ids. There were a small number of null values, and it was found that all of these null values likely belonged to a single user (which is how they are stored using the function below).

```
In [73]: df.article_id.value_counts().head()
```

```
Out[73]: 1429.0    937
1330.0    927
1431.0    671
1427.0    643
1364.0    627
Name: article_id, dtype: int64
```

```
In [74]: most_viewed_article_id = "1429.0"# The most viewed article in the dataset as a string u
max_views = 937# The most viewed article in the dataset was viewed how many times?
```

```
In [75]: ## No need to change the code here - this will be helpful for later parts of the noteb
# Run this cell to map the user email to a user_id column and remove the email column
```

```
def email_mapper():
    coded_dict = dict()
    cter = 1
    email_encoded = []

    for val in df['email']:
        if val not in coded_dict:
            coded_dict[val] = cter
            cter+=1

        email_encoded.append(coded_dict[val])
    return email_encoded
```

```
email_encoded = email_mapper()
del df['email']
df['user_id'] = email_encoded
```

```
# show header
df.head()
```

```
Out[75]:
```

	article_id	title	user_id
0	1430.0	using pixiedust for fast, flexible, and easier...	1
1	1314.0	healthcare python streaming application demo	2
2	1429.0	use deep learning for image classification	3

3	1338.0	ml optimization using cognitive assistant	4
4	1276.0	deploy your python model as a restful api	5

```
In [76]: ## If you stored all your results in the variable names above,
        ## you shouldn't need to change anything in this cell
```

```
sol_1_dict = {
    '50% of individuals have ____ or fewer interactions.': median_val,
    'The total number of user-article interactions in the dataset is ____.': user_a
    'The maximum number of user-article interactions by any 1 user is ____.': max_v
    'The most viewed article in the dataset was viewed ____ times.': max_views,
    'The article_id of the most viewed article is ____.': most_viewed_article_id,
    'The number of unique articles that have at least 1 rating ____.': unique_artic
    'The number of unique users in the dataset is ____.': unique_users,
    'The number of unique articles on the IBM platform': total_articles
}

# Test your dictionary against the solution
t.sol_1_test(sol_1_dict)
```

It looks like you have everything right here! Nice job!

```
In [16]: df.groupby('article_id')['user_id'].count().reset_index(name='count').sort_values(['cou
```

```
Out[16]:
```

	article_id	count
699	1429.0	937
625	1330.0	927
701	1431.0	671
697	1427.0	643
652	1364.0	627
614	1314.0	614
600	1293.0	572
526	1170.0	565
518	1162.0	512
608	1304.0	483
706	1436.0	481
583	1271.0	473
671	1398.0	465
22	43.0	460
641	1351.0	457
666	1393.0	455
540	1185.0	442
516	1160.0	433
642	1354.0	426
656	1368.0	418
609	1305.0	413
633	1338.0	382
631	1336.0	379

521	1165.0	372
528	1172.0	363
76	151.0	352
586	1276.0	347
702	1432.0	340
700	1430.0	336
438	1052.0	330
..
224	499.0	2
385	940.0	2
637	1346.0	2
523	1167.0	2
548	1195.0	2
634	1340.0	2
403	972.0	2
550	1197.0	2
630	1335.0	2
653	1365.0	2
629	1334.0	2
358	870.0	2
310	724.0	1
567	1233.0	1
458	1072.0	1
405	974.0	1
570	1237.0	1
189	417.0	1
554	1202.0	1
289	675.0	1
282	662.0	1
553	1200.0	1
472	1092.0	1
277	653.0	1
636	1344.0	1
478	1113.0	1
481	1119.0	1
409	984.0	1
488	1127.0	1
581	1266.0	1

[714 rows x 2 columns]

1.1.2 Part II: Rank-Based Recommendations

Unlike in the earlier lessons, we don't actually have ratings for whether a user liked an article or not. We only know that a user has interacted with an article. In these cases, the popularity of an article can really only be based on how often an article was interacted with.

1. Fill in the function below to return the **n** top articles ordered with most interactions as the top. Test your function using the tests below.

```

In [77]: def get_top_articles(n, df=df):
        '''
        INPUT:
        n - (int) the number of top articles to return
        df - (pandas dataframe) df as defined at the top of the notebook

        OUTPUT:
        top_articles - (list) A list of the top 'n' article titles

        '''
        df_sorted = df.groupby('article_id').count().sort_values('title', ascending=False)
        top_ids = df_sorted.index.values[:n]
        top_articles = [df.title[df.article_id == id].iloc[0] for id in top_ids]

        return list(top_articles)

def get_top_article_ids(n, df=df):
    '''
    INPUT:
    n - (int) the number of top articles to return
    df - (pandas dataframe) df as defined at the top of the notebook

    OUTPUT:
    top_articles - (list) A list of the top 'n' article titles

    '''
    # Your code here
    top_articles_df = df.groupby('article_id')['user_id'].count().reset_index(name='count')
    top_articles = list(top_articles_df['article_id'][:n])
    return top_articles # Return the top article ids

In [78]: print(get_top_articles(10))

['use deep learning for image classification', 'insights from new york car accident reports', 'v

In [79]: print(get_top_article_ids(10))

[1429.0, 1330.0, 1431.0, 1427.0, 1364.0, 1314.0, 1293.0, 1170.0, 1162.0, 1304.0]

In [80]: # Test your function by returning the top 5, 10, and 20 articles
top_5 = get_top_articles(5)
top_10 = get_top_articles(10)
top_20 = get_top_articles(20)

# Test each of your three lists from above
t.sol_2_test(get_top_articles)

```


Your top_5 looks like the solution list! Nice job.
Your top_10 looks like the solution list! Nice job.
Your top_20 looks like the solution list! Nice job.

1.1.3 Part III: User-User Based Collaborative Filtering

1. Use the function below to reformat the **df** dataframe to be shaped with users as the rows and articles as the columns.

- Each **user** should only appear in each **row** once.
- Each **article** should only show up in one **column**.
- **If a user has interacted with an article, then place a 1 where the user-row meets for that article-column.** It does not matter how many times a user has interacted with the article, all entries where a user has interacted with an article should be a 1.
- **If a user has not interacted with an item, then place a zero where the user-row meets for that article-column.**

Use the tests to make sure the basic structure of your matrix matches what is expected by the solution.

```
In [83]: def create_user_item_matrix(df):
        '''
        INPUT:
        df - pandas dataframe with article_id, title, user_id columns

        OUTPUT:
        user_item - user item matrix

        Description:
        Return a matrix with user ids as rows and article ids on the columns with 1 values
        an article and a 0 otherwise
        '''
        # Fill in the function here
        user_item = df.groupby(["user_id", "article_id"])["title"].count().unstack()

        user_item = user_item.fillna(value=0)

        user_item[user_item > 0] = 1

        return user_item
user_item = create_user_item_matrix(df)

In [84]: ## Tests: You should just need to run this cell. Don't change the code.
assert user_item.shape[0] == 5149, "Oops! The number of users in the user-article matrix is not 5149"
assert user_item.shape[1] == 714, "Oops! The number of articles in the user-article matrix is not 714"
assert user_item.sum(axis=1)[1] == 36, "Oops! The number of articles seen by user 1 does not equal 36"
print("You have passed our quick tests! Please proceed!")
```

You have passed our quick tests! Please proceed!

2. Complete the function below which should take a `user_id` and provide an ordered list of the most similar users to that user (from most similar to least similar). The returned result should not contain the provided `user_id`, as we know that each user is similar to him/herself. Because the results for each user here are binary, it (perhaps) makes sense to compute similarity as the dot product of two users.

Use the tests to test your function.

```
In [85]: def find_similar_users(user_id, user_item=user_item):
        """
        INPUT:
        user_id - (int) a user_id
        user_item - (pandas dataframe) matrix of users by articles:
                    1's when a user has interacted with an article, 0 otherwise

        OUTPUT:
        similar_users - (list) an ordered list where the closest users (largest dot product)
                        are listed first

        Description:
        Computes the similarity of every pair of users based on the dot product
        Returns an ordered

        """
        # compute similarity of each user to the provided user
        similarity = user_item[user_item.index == user_id].dot(user_item.T)
        # sort by similarity
        # create list of just the ids

        most_similar_users = similarity.sort_values(user_id, axis=1, ascending=False).columns

        # remove the own user's id
        most_similar_users.remove(user_id)
        return most_similar_users # return a list of the users in order from most to least

In [86]: # Do a spot check of your function
        print("The 10 most similar users to user 1 are: {}".format(find_similar_users(1)[:10]))
        print("The 5 most similar users to user 3933 are: {}".format(find_similar_users(3933)[:5]))
        print("The 3 most similar users to user 46 are: {}".format(find_similar_users(46)[:3]))

The 10 most similar users to user 1 are: [3933, 23, 3782, 203, 4459, 3870, 131, 4201, 46, 5041]
The 5 most similar users to user 3933 are: [1, 23, 3782, 203, 4459]
The 3 most similar users to user 46 are: [4201, 3782, 23]
```

```
In [87]: user_item.loc[1][user_item.loc[1] == 1].reset_index()['article_id'].values
```

```
Out[87]: array([[ 43.,  109.,  151.,  268.,  310.,  329.,  346.,  390.,
        494.,  525.,  585.,  626.,  668.,  732.,  768.,  910.,
        968.,  981., 1052., 1170., 1183., 1185., 1232., 1293.,
       1305., 1363., 1368., 1391., 1400., 1406., 1427., 1429.,
       1430., 1431., 1436., 1439.]])
```

3. Now that you have a function that provides the most similar users to each user, you will want to use these users to find articles you can recommend. Complete the functions below to return the articles you would recommend to each user.

```
In [88]: def get_article_names(article_ids, df=df):
        '''
        INPUT:
        article_ids - (list) a list of article ids
        df - (pandas dataframe) df as defined at the top of the notebook

        OUTPUT:
        article_names - (list) a list of article names associated with the list of article
                        (this is identified by the title column)
        '''
        # Your code here
        article_names = df[df['article_id'].isin(article_ids)]['title'].drop_duplicates().v
        return article_names # Return the article names associated with list of article ids


def get_user_articles(user_id, user_item=user_item):
    '''
    INPUT:
    user_id - (int) a user id
    user_item - (pandas dataframe) matrix of users by articles:
                1's when a user has interacted with an article, 0 otherwise

    OUTPUT:
    article_ids - (list) a list of the article ids seen by the user
    article_names - (list) a list of article names associated with the list of article
                    (this is identified by the doc_full_name column in df_content)

    Description:
    Provides a list of the article_ids and article titles that have been seen by a user
    '''
    # Your code here
    article_ids = user_item.loc[user_id][user_item.loc[user_id]!=0].reset_index()['arti
    article_ids = [str(article_id) for article_id in article_ids]
    article_names = get_article_names(article_ids)
    return article_ids, article_names # return the ids and names


def user_user_recs(user_id, m=10):
```

```

'''
INPUT:
user_id - (int) a user id
m - (int) the number of recommendations you want for the user

OUTPUT:
recs - (list) a list of recommendations for the user

Description:
Loops through the users based on closeness to the input user_id
For each user - finds articles the user hasn't seen before and provides them as recs
Does this until m recommendations are found

Notes:
Users who are the same closeness are chosen arbitrarily as the 'next' user

For the user where the number of recommended articles starts below m
and ends exceeding m, the last items are chosen arbitrarily

'''
# Your code here
most_similar_users = find_similar_users(user_id, user_item)

recs = []

user_seen_articles, _ = get_user_articles(user_id, user_item)
user_seen_articles = set(user_seen_articles)

for user in most_similar_users:
    if len(recs) < m:
        article_ids = set(get_user_articles(user, user_item)[0])
        recommended_ids = article_ids.difference(user_seen_articles)
        req_recs = m - len(recs)
        recommended_ids = list(recommended_ids)[:req_recs]
        recs.extend(recommended_ids)
        user_seen_articles = user_seen_articles.union(recommended_ids)

return recs

```

In [89]: # Check Results

```
get_article_names(user_user_recs(1, 10)) # Return 10 recommendations for user 1
```

```

Out[89]: ['data science for real-time streaming analytics',
'discover hidden facebook usage insights',
'analyze accident reports on amazon emr spark',
'better together: spss and data science experience',
'520    using notebooks with pixiedust for fast, flexi...\nName: title, dtype: object',
'learn tensorflow and deep learning together and now!',

```



```

for user in user_item.index:
    if user == user_id:
        continue
    neighbors_df.loc[user] = [user, np.dot(user_item.loc[user_id, :], user_item.loc[
        df[df['user_id']==user]['article_id'].count())

neighbors_df.sort_values(by=['similarity', 'num_interactions'], ascending=False, in

return neighbors_df # Return the dataframe specified in the doc_string


def user_user_recs_part2(user_id, m=10):
    '''
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user by article id
    rec_names - (list) a list of recommendations for the user by article title

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them as recs
    Does this until m recommendations are found

    Notes:
    * Choose the users that have the most total article interactions
    before choosing those with fewer article interactions.

    * Choose articles with the articles with the most total interactions
    before choosing those with fewer total interactions.

    '''
    # Your code here
    recs = []

    neighbors_df = get_top_sorted_users(user_id)

    the_user_articles, the_article_names = get_user_articles(user_id)
    for user in neighbors_df['neighbor_id']:
        article_ids, article_names = get_user_articles(user)
        for id in article_ids:
            if id not in the_user_articles:
                recs.append(id)
            if len(recs) >= m:
                break
        if len(recs) >= m:

```

```

        break

    if len(recs) < m:
        for id in [str(id) for id in get_top_article_ids(100)]:
            if id not in the_user_articles:
                recs.append(id)
            if len(recs) >= m:
                break

    rec_names = get_article_names(recs)

    return recs, rec_names

```

```

In [92]: # Quick spot check - don't change this code - just use it to test your functions
rec_ids, rec_names = user_user_recs_part2(20, 10)
print("The top 10 recommendations for user 20 are the following article ids:")
print(rec_ids)
print()
print("The top 10 recommendations for user 20 are the following article names:")
print(rec_names)

```

The top 10 recommendations for user 20 are the following article ids:

```
['12.0', '109.0', '125.0', '142.0', '164.0', '205.0', '302.0', '336.0', '362.0', '465.0']
```

The top 10 recommendations for user 20 are the following article names:

```
['timeseries data analysis of iot events by using jupyter notebook', 'dsx: hybrid mode', 'accele
```

5. Use your functions from above to correctly fill in the solutions to the dictionary below. Then test your dictionary against the solution. Provide the code you need to answer each following the comments below.

```

In [93]: ### Tests with a dictionary of results
user1_most_sim = get_top_sorted_users(1).loc[0].neighbor_id# Find the user that is most
user131_10th_sim = get_top_sorted_users(131).loc[9].neighbor_id# Find the 10th most sim

```

KeyError

Traceback (most recent call last)

```

/opt/conda/lib/python3.6/site-packages/pandas/core/indexing.py in _validate_key(self, key)
1789         if not ax.contains(key):
-> 1790             error()
1791         except TypeError as e:

/opt/conda/lib/python3.6/site-packages/pandas/core/indexing.py in error()
1784         .format(key=key,

```

```
-> 1785                 axis=self.obj._get_axis_name(axis)))
    1786
```

```
KeyError: 'the label [0] is not in the [index]'
```

During handling of the above exception, another exception occurred:

```
KeyError                                Traceback (most recent call last)
```

```
<ipython-input-93-d85a71facf62> in <module>()
    1 ### Tests with a dictionary of results
----> 2 user1_most_sim = get_top_sorted_users(1).loc[0].neighbor_id# Find the user that is m
    3 user131_10th_sim = get_top_sorted_users(131).loc[9].neighbor_id# Find the 10th most
```

```
/opt/conda/lib/python3.6/site-packages/pandas/core/indexing.py in __getitem__(self, key)
1476
1477         maybe_callable = com._apply_if_callable(key, self.obj)
-> 1478         return self._getitem_axis(maybe_callable, axis=axis)
1479
1480     def _is_scalar_access(self, key):
```

```
/opt/conda/lib/python3.6/site-packages/pandas/core/indexing.py in _getitem_axis(self, key, axis)
1909
1910         # fall thru to straight lookup
-> 1911         self._validate_key(key, axis)
1912         return self._get_label(key, axis=axis)
1913
```

```
/opt/conda/lib/python3.6/site-packages/pandas/core/indexing.py in _validate_key(self, key, axis)
1796             raise
1797         except:
-> 1798             error()
1799
1800     def _is_scalar_access(self, key):
```

```
/opt/conda/lib/python3.6/site-packages/pandas/core/indexing.py in error()
1783         raise KeyError(u"the label [{key}] is not in the [{axis}]"
1784                        .format(key=key,
-> 1785                               axis=self.obj._get_axis_name(axis)))
1786
1787     try:
```



```
KeyError: 'the label [0] is not in the [index]'
```

```
In [ ]: ## Dictionary Test Here
sol_5_dict = {
    'The user that is most similar to user 1.': user1_most_sim,
    'The user that is the 10th most similar to user 131': user131_10th_sim,
}

t.sol_5_test(sol_5_dict)
```

6. If we were given a new user, which of the above functions would you be able to use to make recommendations? Explain. Can you think of a better way we might make recommendations? Use the cell below to explain a better method for new users.

ANSWER: `get_top_article_ids` would be a better way for us to make recommendations because the user didn't view any articles before and we don't have any information about user

7. Using your existing functions, provide the top 10 recommended articles you would provide for the a new user below. You can test your function against our thoughts to make sure we are all on the same page with how we might make a recommendation.

```
In [ ]: new_user = '0.0'

# What would your recommendations be for this new user '0.0'? As a new user, they have
# Provide a list of the top 10 article ids you would give to
top_article_ids = get_top_article_ids(10)
new_user_recs = [str(x) for x in top_article_ids]
new_user_recs

In [ ]: assert set(new_user_recs) == set(['1314.0', '1429.0', '1293.0', '1427.0', '1162.0', '1364.0',

print("That's right! Nice job!")
```

1.1.4 Part IV: Content Based Recommendations (EXTRA - NOT REQUIRED)

Another method we might use to make recommendations is to perform a ranking of the highest ranked articles associated with some term. You might consider content to be the **doc_body**, **doc_description**, or **doc_full_name**. There isn't one way to create a content based recommendation, especially considering that each of these columns hold content related information.

1. Use the function body below to create a content based recommender. Since there isn't one right answer for this recommendation tactic, no test functions are provided. Feel free to change the function inputs if you decide you want to try a method that requires more input values. The input values are currently set with one idea in mind that you may use to make content based recommendations. One additional idea is that you might want to choose the most popular recommendations that meet your 'content criteria', but again, there is a lot of flexibility in how you might make these recommendations.

1.1.5 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

```
In [94]: def make_content_recs():  
        '''  
        INPUT:  
  
        OUTPUT:  
  
        '''
```

2. Now that you have put together your content-based recommendation system, use the cell below to write a summary explaining how your content based recommender works. Do you see any possible improvements that could be made to your function? Is there anything novel about your content based recommender?

1.1.6 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

Write an explanation of your content based recommendation system here.

3. Use your content-recommendation system to make recommendations for the below scenarios based on the comments. Again no tests are provided here, because there isn't one right answer that could be used to find these content based recommendations.

1.1.7 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

```
In [ ]: # make recommendations for a brand new user  
  
        # make a recommendations for a user who only has interacted with article id '1427.0'
```

1.1.8 Part V: Matrix Factorization

In this part of the notebook, you will build use matrix factorization to make article recommendations to the users on the IBM Watson Studio platform.

1. You should have already created a **user_item** matrix above in **question 1** of **Part III** above. This first question here will just require that you run the cells to get things set up for the rest of **Part V** of the notebook.

```
In [95]: # Load the matrix here  
user_item_matrix = pd.read_pickle('user_item_matrix.p')
```

```
In [96]: # quick look at the matrix  
user_item_matrix.head()
```

```
Out[96]: article_id  0.0  100.0  1000.0  1004.0  1006.0  1008.0  101.0  1014.0  1015.0  \  
user_id  
1          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
```

2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

article_id	1016.0	...	977.0	98.0	981.0	984.0	985.0	986.0	990.0	\
user_id		...								
1	0.0	...	0.0	0.0	1.0	0.0	0.0	0.0	0.0	
2	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	...	1.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
5	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

article_id	993.0	996.0	997.0
user_id			
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0
5	0.0	0.0	0.0

[5 rows x 714 columns]

2. In this situation, you can use Singular Value Decomposition from [numpy](#) on the user-item matrix. Use the cell to perform SVD, and explain why this is different than in the lesson.

In [97]: *# Perform SVD on the User-Item Matrix Here*

```
u, s, vt = np.linalg.svd(user_item_matrix) # use the built in to get the three matrices
```

ANSWER: SVD can only be used when we have no missing values in the collaborative filtering matrix. In this exercise we have replaced all Nan values with 0's otherwise we would have needed to use FunkSVD.

3. Now for the tricky part, how do we choose the number of latent features to use? Running the below cell, you can see that as the number of latent features increases, we obtain a lower error rate on making predictions for the 1 and 0 values in the user-item matrix. Run the cell below to get an idea of how the accuracy improves as we increase the number of latent features.

```
In [98]: num_latent_feats = np.arange(10, 700+10, 20)
sum_errs = []
```

```
for k in num_latent_feats:
    # restructure with k latent features
    s_new, u_new, vt_new = np.diag(s[:k]), u[:, :k], vt[:k, :]

    # take dot product
    user_item_est = np.around(np.dot(np.dot(u_new, s_new), vt_new))

    # compute error for each prediction to actual value
```


```

diffs = np.subtract(user_item_matrix, user_item_est)

# total errors and keep track of them
err = np.sum(np.sum(np.abs(diffs)))
sum_errs.append(err)

plt.plot(num_latent_feats, 1 - np.array(sum_errs)/df.shape[0]);
plt.xlabel('Number of Latent Features');
plt.ylabel('Accuracy');
plt.title('Accuracy vs. Number of Latent Features');

```



4. From the above, we can't really be sure how many features to use, because simply having a better way to predict the 1's and 0's of the matrix doesn't exactly give us an indication of if we are able to make good recommendations. Instead, we might split our dataset into a training and test set of data, as shown in the cell below.

Use the code from question 3 to understand the impact on accuracy of the training and test sets of data with different numbers of latent features. Using the split below:

- How many users can we make predictions for in the test set?
- How many users are we not able to make predictions for because of the cold start problem?
- How many articles can we make predictions for in the test set?
- How many articles are we not able to make predictions for because of the cold start problem?

```

In [99]: df_train = df.head(40000)
         df_test = df.tail(5993)

def create_test_and_train_user_item(df_train, df_test):
    '''
    INPUT:
    df_train - training dataframe
    df_test - test dataframe

    OUTPUT:
    user_item_train - a user-item matrix of the training dataframe
                      (unique users for each row and unique articles for each column)
    user_item_test - a user-item matrix of the testing dataframe
                     (unique users for each row and unique articles for each column)
    test_idx - all of the test user ids
    test_arts - all of the test article ids

    '''
    # Your code here
    user_item_train = create_user_item_matrix(df_train)
    user_item_test = create_user_item_matrix(df_test)

    train_idx = set(user_item_test.index.tolist())
    test_idx = set(user_item_test.index.tolist())
    duplicate_idx = list(train_idx.intersection(test_idx))

    train_arts = set(user_item_train.columns.tolist())
    test_arts = set(user_item_test.columns.tolist())
    duplicate_arts = list(train_arts.intersection(test_arts))

    user_item_test = user_item_test.loc[duplicate_idx, duplicate_arts]
    return user_item_train, user_item_test, test_idx, test_arts

user_item_train, user_item_test, test_idx, test_arts = create_test_and_train_user_item(

In [100]: #How many users can we make predictions for in the test set?
          unique_test_pred = set(user_item_test.index.tolist())&set(user_item_train.index.tolist())
          len(unique_test_pred)

Out[100]: 20

In [101]: #How many users in the test set are we not able to make predictions for because of the
          cold_start_pred = set(user_item_test.index.tolist()) - set(user_item_train.index.tolist())
          len(cold_start_pred)

Out[101]: 662

In [102]: #How many articles can we make predictions for in the test set?
          unique_art_predication = set(user_item_test.columns.tolist())&set(user_item_train.columns.tolist())
          len(unique_art_predication)

```

Out[102]: 574

```
In [103]: #How many articles in the test set are we not able to make predictions
unique_art_predication = set(user_item_test.columns.tolist())-set(user_item_train.columns.tolist())
len(unique_art_predication)
```

Out[103]: 0

```
In [104]: # Replace the values in the dictionary below
a = 662
b = 574
c = 20
d = 0

sol_4_dict = {
    'How many users can we make predictions for in the test set?': c,
    'How many users in the test set are we not able to make predictions for because of': d,
    'How many movies can we make predictions for in the test set?': b,
    'How many movies in the test set are we not able to make predictions for because of': a
}

t.sol_4_test(sol_4_dict)
```

Awesome job! That's right! All of the test movies are in the training data, but there are only

```
In [105]: user_item_test.shape
```

Out[105]: (682, 574)

5. Now use the **user_item_train** dataset from above to find U, S, and V transpose using SVD. Then find the subset of rows in the **user_item_test** dataset that you can predict using this matrix decomposition with different numbers of latent features to see how many features makes sense to keep based on the accuracy on the test data. This will require combining what was done in questions 2 - 4.

Use the cells below to explore how well SVD works towards making predictions for recommendations on the test data.

```
In [106]: # fit SVD on the user_item_train matrix
u_train, s_train, vt_train = np.linalg.svd(np.array(user_item_train, dtype='int'), full=True)
```

```
In [107]: # Use these cells to see how well you can use the training
# decomposition to predict on test data
row_idxes = user_item_train.index.isin(test_idxes)
cols_idxes = user_item_train.columns.isin(test_arts)
u_test = u_train[row_idxes, :]
vt_test = vt_train[:, cols_idxes]
```

```

In [108]: num_latent_feats = np.arange(10,700+10,20)
          sum_errs_train = []
          sum_errs_test = []
          all_errs = []
          user_item_test = user_item_test.loc[unique_test_pred,: ]

          for k in num_latent_feats:
              # restructure with k latent features
              s_train_lat, u_train_lat, vt_train_lat = np.diag(s_train[:k]), u_train[:, :k], vt_train[:, :k]

              u_test_lat, vt_test_lat = u_test[:, :k], vt_test[:, :k]


              # take dot product
              user_item_train_preds = np.around(np.dot(np.dot(u_train_lat, s_train_lat), vt_train_lat))
              user_item_test_preds = np.around(np.dot(np.dot(u_test_lat, s_train_lat), vt_test_lat))
              all_errs.append(1 - ((np.sum(user_item_test_preds)+np.sum(np.sum(user_item_test))))

              diffs_train = np.subtract(user_item_train, user_item_train_preds)
              diffs_test = np.subtract(user_item_test, user_item_test_preds)

              # total errors and keep track of them
              err_train = np.sum(np.sum(np.abs(diffs_train)))
              err_test = np.sum(np.sum(np.abs(diffs_test)))
              sum_errs_train.append(err_train)
              sum_errs_test.append(err_test)

In [109]: plt.plot(num_latent_feats, 1 - np.array(sum_errs_train)/(user_item_train.shape[0]*user_item_train.shape[1]))
          plt.plot(num_latent_feats, 1 - np.array(sum_errs_test)/(user_item_test.shape[0]*user_item_test.shape[1]))
          plt.plot(num_latent_feats, all_errs, label='All Data');
          plt.xlabel('Number of Latent Features');
          plt.ylabel('Accuracy');
          plt.title('Accuracy vs. Number of Latent Features');
          plt.legend();

```



6. Use the cell below to comment on the results you found in the previous question. Given the circumstances of your results, discuss what you might do to determine if the recommendations you make with any of the above recommendation systems are an improvement to how users currently find articles?

Your response here. Increasing the latent features causes overfitting problem. Thus we can see from the above graph, the accuracy becomes worse when the number of latent features increases. Since the common users between the train and test set are too few, other recommendation methods may be used to improve our recommendation, like collaborative filtering or content based recommendation

We could use A/B testing to test how well our recommendation engine is working in practice to further engage users. We separate two groups of user, one uses our recommendation engine and another uses random recommendation. we compare the hit rate of the recommendation articles to measure if our recommendation engine boost up the view count. If it is significant, we can conclude that our recommendation engine works well.

Extras Using your workbook, you could now save your recommendations for each user, develop a class to make new predictions and update your results, and make a flask app to deploy your results. These tasks are beyond what is required for this project. However, from what you learned in the lessons, you certainly capable of taking these tasks on to improve upon your work here!

1.2 Conclusion

Congratulations! You have reached the end of the Recommendations with IBM project!

Tip: Once you are satisfied with your work here, check over your report to make sure that it satisfies all the areas of the [rubric](#). You should also probably remove all of the "Tips" like this one so that the presentation is as polished as possible.

1.3 Directions to Submit

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File > Download as** sub-menu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

```
In [110]: from subprocess import call
          call(['python', '-m', 'nbconvert', 'Recommendations_with_IBM.ipynb'])
```

```
Out[110]: 0
```

```
In [ ]:
```